



Cloud Computing

Introduction to Apache Spark

Seyyed Ahmad Javadi

sajavadi@aut.ac.ir

Fall 2021



Apache Spark 101

Lance Co Ting Keh

Senior Software Engineer, Machine Learning @ Box

Outline



- I. About me
- II. Distributed Computing at a High Level
- III. Disk versus Memory based Systems
- IV. Spark Core
 - I. Brief background
 - II. Benchmarks and Comparisons
 - III. What is an RDD
 - IV. RDD Actions and Transformations
 - V. Caching and Serialization
 - VI. Anatomy of a Program
 - VII. The Spark Family

Why Distributed Computing?



Divide and Conquer

Problem Single machine cannot complete the computation at hand

Solution Parallelize the job and distribute work among a network of machines



Issues Arise in Distributed Computing



View the world from the eyes of a single worker

- How do I **distribute** an algorithm?
- How do I **partition** my dataset?
- How do I maintain a **single consistent view** of a shared state?
- How do I recover from **machine failures**?
- How do I **allocate cluster resources**?
-

Finding majority element in a single machine

Think distributed



List(20, 18, 20, 18, 20)

Finding majority element in a distributed dataset



Think distributed

List(1, 18, 1, 18, 1)

List(2, 18, 2, 18, 2)

List(3, 18, 3, 18, 3)

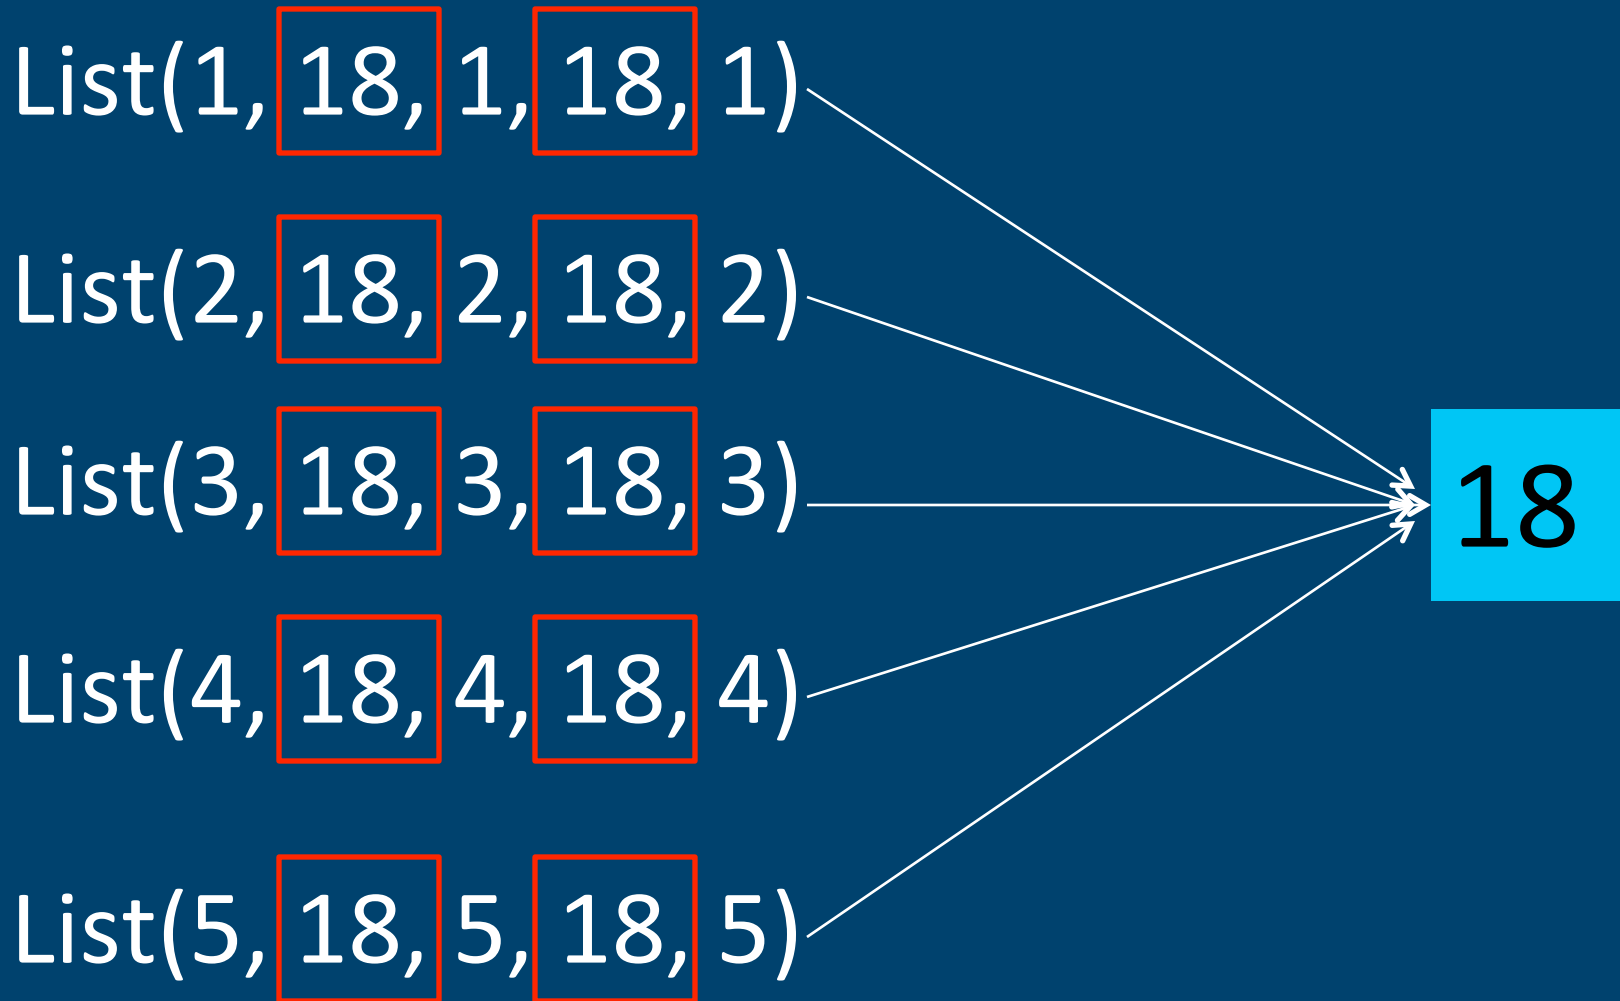
List(4, 18, 4, 18, 4)

List(5, 18, 5, 18, 5)

Finding majority element in a distributed dataset



Think distributed



Disk Based vs Memory Based Frameworks



Acyclic data flow

- Disk Based Frameworks

- Persists intermediate results to disk
- Data is reloaded from disk with every query
- Easy failure recovery
- Best for ETL like work-loads
- Examples: Hadoop, Dryad

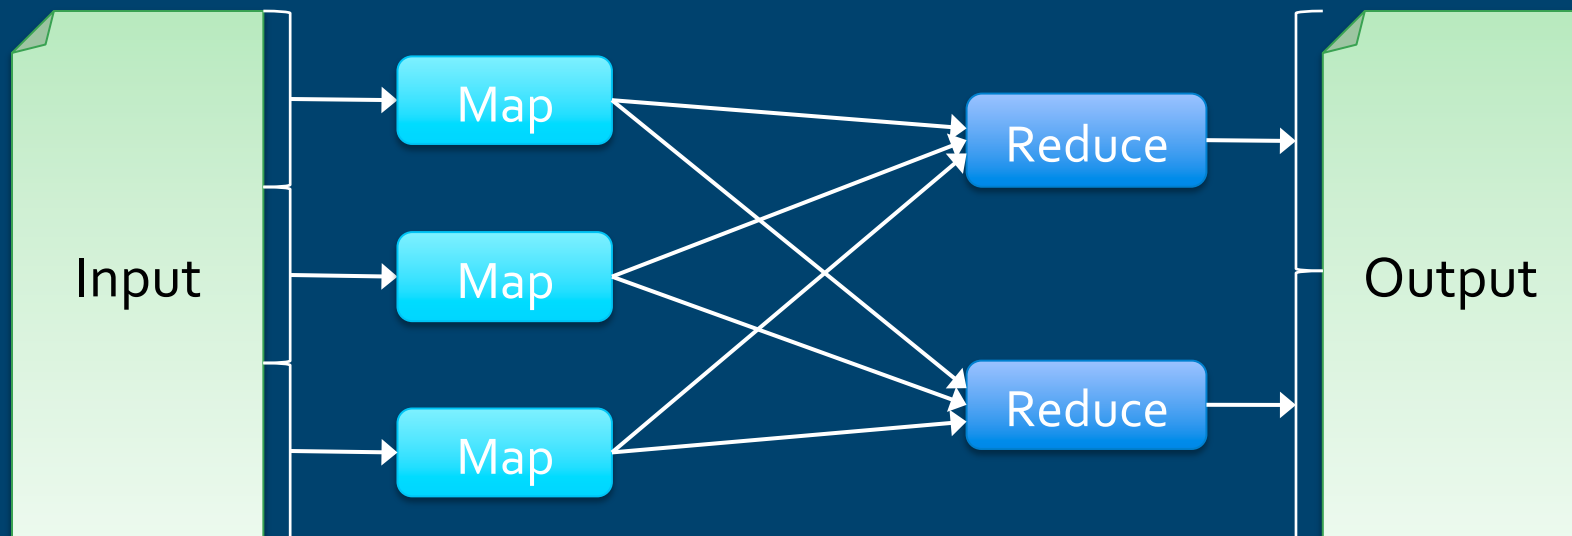


Image courtesy of Matei Zaharia, Introduction to Spark

Disk Based vs Memory Based Frameworks

Reuse working data set in memory

- Memory Based Frameworks
 - Circumvents heavy cost of I/O by keeping intermediate results in memory
 - Sensitive to availability of memory
 - Remembers operations applied to dataset
 - Best for iterative workloads
 - Examples: Spark, Flink




Image courtesy of Matei Zaharia, Introduction to Spark



The rest of the talk

I. Spark Core

- I. Brief background
- II. Benchmarks and Comparisons
- III. What is an RDD
- IV. RDD Actions and Transformations
- V. Spark Cluster
- VI. Anatomy of a Program
- VII. The Spark Family



Spark Background

Arose from an academic setting

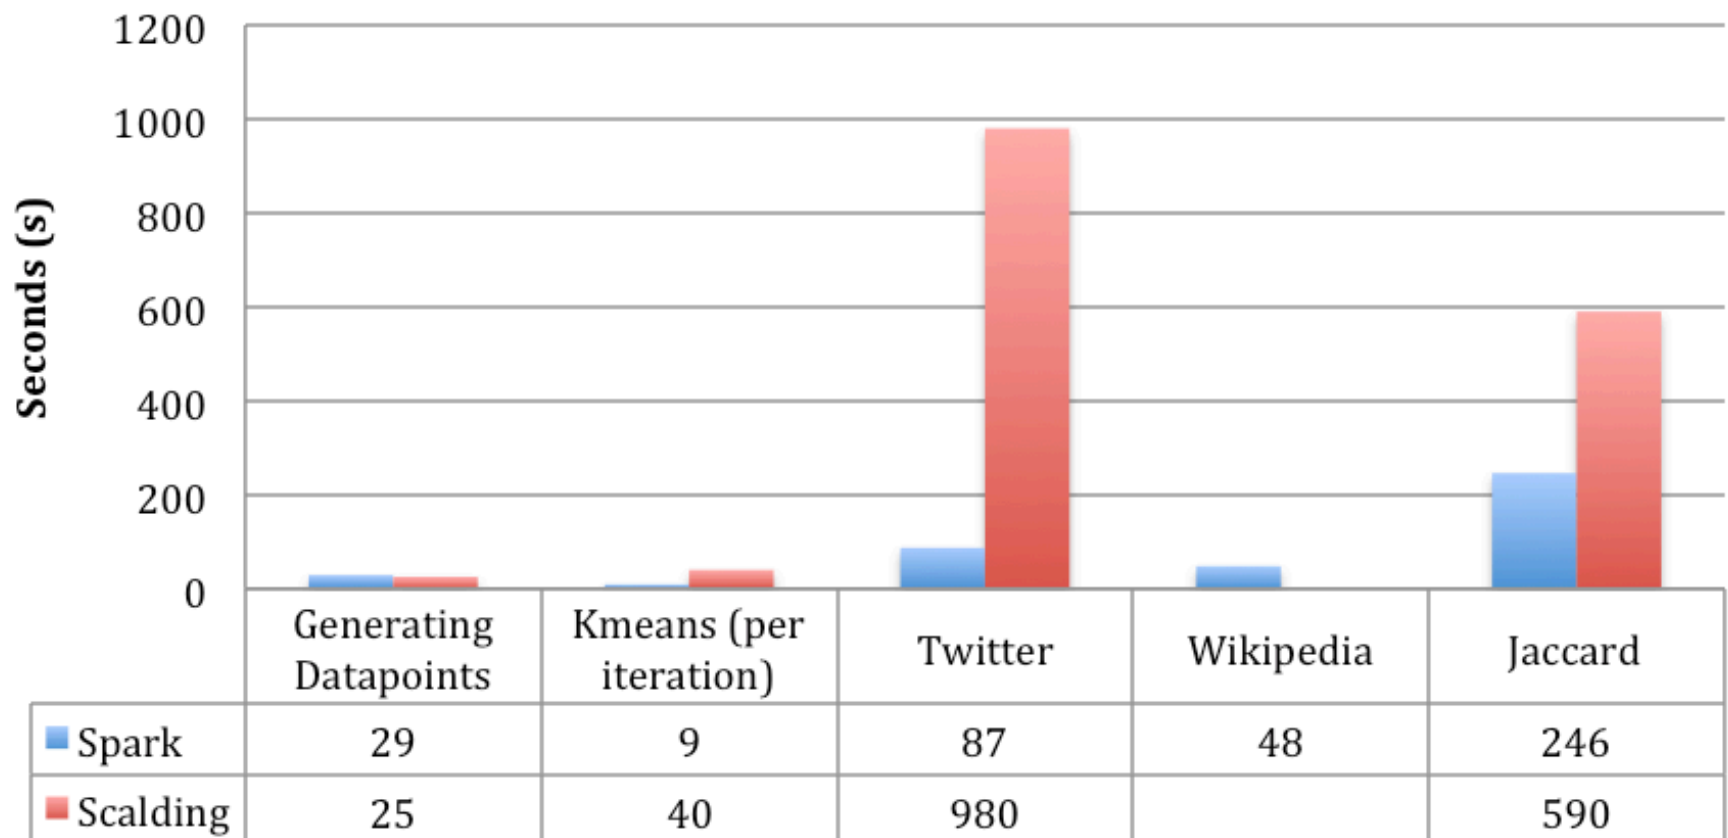
- Amplab UC Berkley
- Project Lead: Dr. Matei Zaharia
- First paper published on RDD's was in 2012
- Open sourced from day one, growing number of contributors
- Released its 1.0 version May 2014. Currently in 1.2.1
- *Databricks* company established to support Spark and all its related technologies. Matei currently sits as its CTO
- Amazon, Alibaba, Baidu, eBay, Groupon, Ooyala, OpenTable, Box, Shopify, TechBase, Yahoo!

Spark versus Scalding (Hadoop)



Clear win for iterative applications

Benchmarks Runtime



Resilient Distributed Datasets (RDDs)



- Main object in Spark's universe
- Think of it as representing the data at that stage in the operation
- Allows for *coarse-grained* transformations (e.g. map, group-by, join)
- Allows for efficient fault recovery using *lineage*
 - Log *one* operation to apply to many elements
 - Recompute lost partitions of dataset on failure
 - No cost if nothing fails

RDD Actions and Transformations



Transformations are realized when an action is called

- Transformations
 - Lazy operations applied on an RDD
 - Creates a new RDD from an existing RDD
 - Allows Spark to perform optimizations
 - e.g. map, filter, flatMap, union, intersection, distinct, reduceByKey, groupByKey
- Actions
 - Returns a value to the driver program after computation
 - e.g. reduce, collect, count, first, take, saveAsFile

RDD Representation



- Simple common interface:
 - Set of partitions
 - Preferred locations for each partition
 - List of parent RDDs
 - Function to compute a partition given parents
 - Optional partitioning info
- Allows capturing wide range of transformations

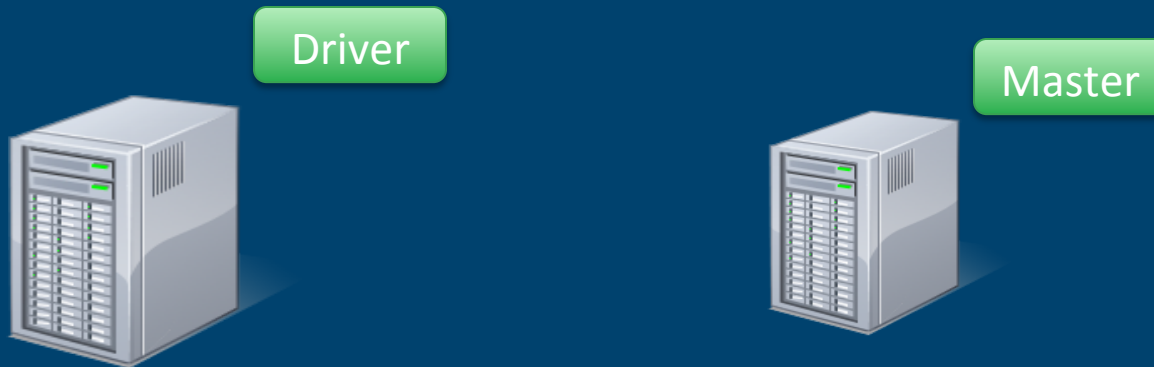
Spark Cluster



Driver

- Entry point of Spark application
- Main Spark application is ran here
- Results of “reduce” operations are aggregated here

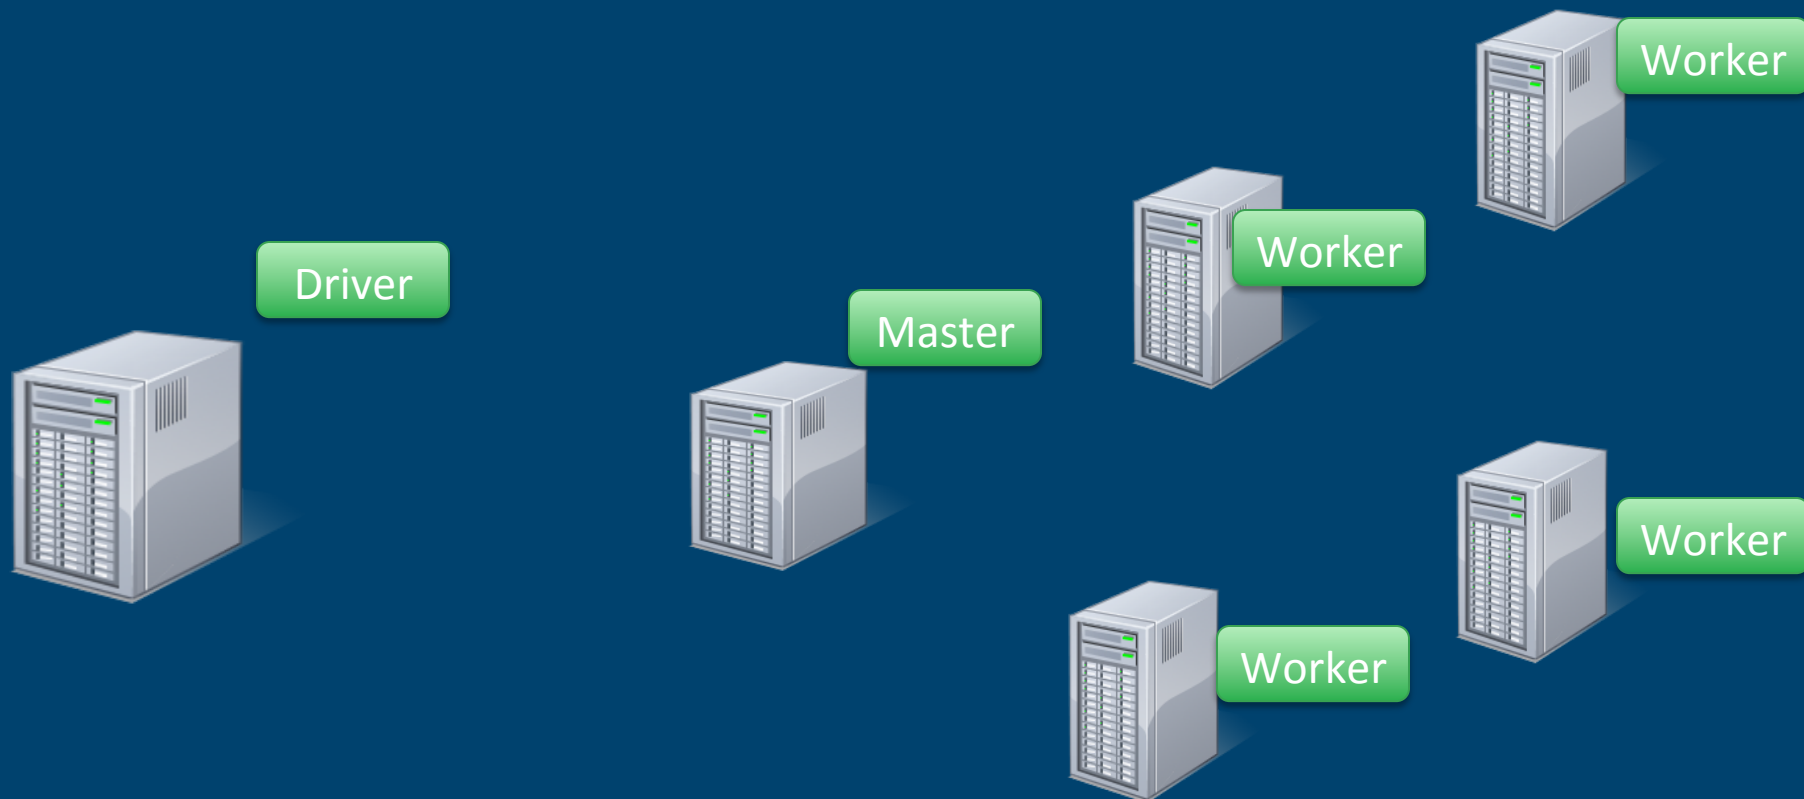
Spark Cluster



Master

- Distributed coordination of Spark workers including:
 - Health checking workers
 - Reassignment of failed tasks
 - Entry point for job and cluster metrics

Spark Cluster



Worker

- Spawns executors to perform tasks on partitions of data

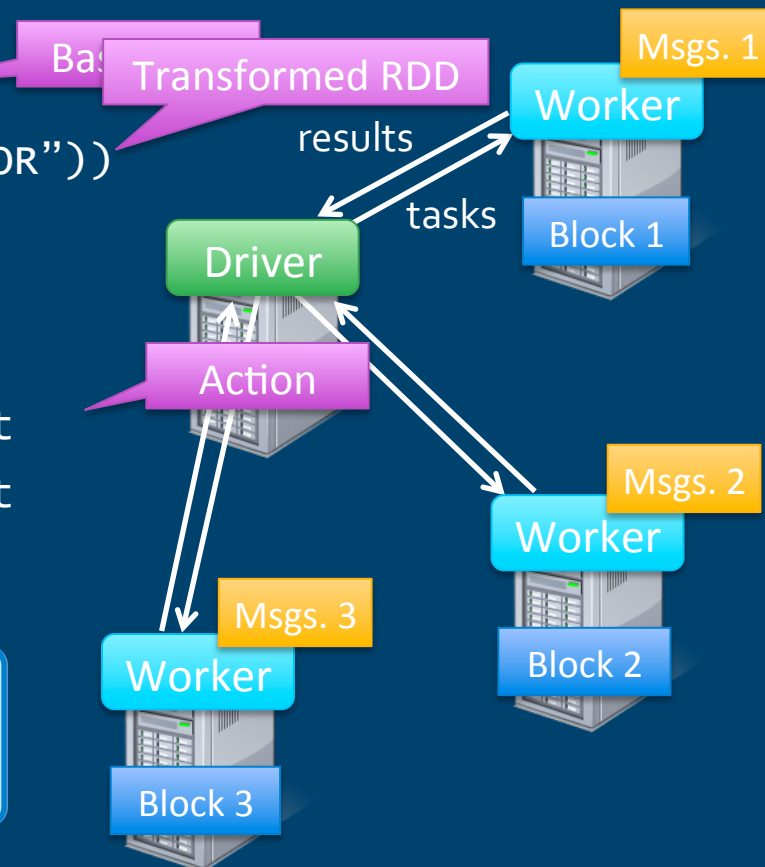
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
messages.persist()
```

```
messages.filter(_.contains("foo")).count
messages.filter(_.contains("bar")).count
. . .
```

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)



The Spark Family



Cheaper by the dozen

- Aside from its performance and API, the diverse tool set available in Spark is the reason for its wide adoption
 1. Spark Streaming
 2. Spark SQL
 3. MLlib
 4. GraphX